GlandersWatch: A Modular Platform for Glanders Risk Prediction, Retrieval-Augmented Explanation and Owner Alerting

Project Report Submitted

in partial fulfillment of the requirements for the award of the

Bachelor of Technology

in

Computer Science Engineering

by

Pedapatruni Shanmuk Vardhan B123089
Shaktimayee Mohanty B123118
Yuvraj Singh B123153

Under the supervision of

Dr. Sangita Dhara



Department of Computer Science and Engineering

International Institute of Information Technology, Bhubaneswar

November 2026

Certificate

This is to certify that the report entitled "GlandersWatch: A Modular Platform for Glanders Risk Prediction, Retrieval-Augmented Explanation and Owner Alerting" submitted by Pedapatruni Shanmuk Vardhan (B123089), Shaktimayee Mohanty (B123118) and Yuvraj Singh (B123153) to International Institute of Information Technology, Bhubaneswar is a record of bonafide project work under my supervision, and the report is submitted for the evaluation of the BTech Degree Final-Semester Project.

Dr. Sangita Dhara
(Supervisor)

Declaration

We certify that

1. The work contained in the report has been done by ourselves under the general supervision of our supervisor.
2. The work has not been submitted to any other Institute for any degree or diploma.
3. We have followed the guidelines provided by the Institute in writing the report.
4. We have conformed to the norms and guidelines given in the Ethical Code of Conduct of the Institute.
5. Whenever we have used materials (data, theoretical analysis, and text) from other sources, we have given due credit to them by citing them in the text of the report and giving their details in the references.
6. Whenever we have quoted written materials from other sources, we have put them under quotation marks and given due credit to the sources by citing them and giving required details in the references.

Pedapatruni Shanmuk Vardhan	**Shaktimayee Mohanty**	**Yuvraj Singh**
B123089	B123118	B123153

Acknowledgments

We would like to express our sincere gratitude to our respected supervisor, Dr. Sangita Dhara, for the invaluable guidance, support and encouragement given throughout this project. Her insights and expertise were instrumental in shaping this work and helped us overcome numerous challenges. We are also thankful to the faculty and staff of the Department of Computer Science and Engineering, International Institute of Information Technology, Bhubaneswar, for providing the resources and a conducive learning environment. Finally, we would like to thank our families and friends for their unwavering support and understanding during the course of this project.

Pedapatruni Shanmuk Vardhan
B123089

Shaktimayee Mohanty
B123118

Yuvraj Singh
B123153

Abstract

Glanders is a bacterial zoonosis of equids caused by *Burkholderia mallei*. It is notifiable, frequently fatal, and remains endemic in parts of Africa, Asia and South America. Indian surveillance between 2019 and 2023 identified 728 seropositive equids across 138 districts, with strong geographic concentration. The published machine learning literature on this disease is confined to laboratory identification of the organism by spectroscopic methods, and no population-level geographic risk model has been reported.

This report presents GlandersWatch, a platform built from three independently deployable components: a calibrated machine learning risk engine, a retrieval-augmented layer that explains each prediction with citations to official guidance, and an alert service that delivers the result to equine owners. The decomposition is the contribution as much as the model is. Each component communicates through a versioned schema rather than shared internals, none hardcodes the disease, and each degrades without failing the others.

The risk engine combines four public datasets not previously applied to glanders into a country-year panel: a text-mined occurrence corpus from Europe PMC, FAOSTAT livestock stocks, World Bank development indicators and NASA POWER reanalysis climatology. The panel contains 3,078 country-year observations across 162 equine-holding countries from 2005 to 2023, with 172 positive labels and 44 engineered features. Six model families are trained under three validation schemes, and a three-stage chained pipeline is evaluated in which animal population is forecast before occurrence is predicted.

On a forward temporal holdout the strongest model reaches a precision-recall AUC of 0.554 against a base rate of 0.056, roughly a ten-fold lift, with an expected calibration error of 0.016. Performance falls to 0.304 under leave-country-out validation, which is the estimate to quote for an unseen country. Chaining produces no measurable gain in discrimination. An ablation shows climate covariates transfer poorly across countries, and the model performs least well in low-income regions. The dominant limitation is that the label reflects documented activity rather than true incidence, so the model largely learns where disease is reported. Reporting effort is the principal confounder and the main target for future work. The explanation and alert components are specified to the level of their interfaces, rule semantics and evaluation protocol, and are not reported as measured results because they are not deployed.

Keywords: Glanders, *Burkholderia mallei*, Risk Prediction, Gradient Boosting, Probability Calibration, SHAP, Retrieval-Augmented Generation, Citation Grounding, Alert Fatigue, Modular System Design, MLOps, Animal Health Surveillance.

Contents

Abstract **4**

1 Introduction **6**

2 Literature Review and Gap Analysis **8**

3 System Architecture **10**

4 Methodology **13**

5 MLOps and Model Serving **15**

6 The RAG Explanation Layer **18**

7 The Alert Service **21**

8 Results and Discussion **24**

9 Conclusion and Future Work **27**

Bibliography **28**

Chapter 1

Introduction

1. Background

Glanders is a contagious and often fatal disease of horses, donkeys and mules caused by *Burkholderia mallei*. The organism has no environmental reservoir, and asymptomatic carriers remain infectious for long periods. Human infection is rare but carries a high case fatality rate and follows occupational contact. The disease is notifiable to the World Organisation for Animal Health and is classified as a potential bioterrorism agent.

The economic weight of the disease sits in low and middle income countries, where equids remain essential to transport, construction, agriculture and brick manufacture. India recorded 5.44 lakh equids in the 2019 Livestock Census, many held by households that depend on them for daily income. The loss of a working animal is a direct loss of livelihood, which makes early warning useful even where mass treatment is not feasible.

Surveillance data show the disease is neither rare nor geographically contained. A national programme screening 158,798 equid sera from 391 districts between 2019 and 2023 found 728 seropositive animals in 138 districts, with roughly sixty per cent of positives in Uttar Pradesh. Earlier surveillance covering 2006 to 2018 reported 1,145 positive equids among 128,404 sera tested. These are persistent, spatially structured figures, not isolated outbreaks.

2. Problem Definition

A district veterinary officer faces a narrow question: is glanders likely to appear in this area in the coming year? Answering it means combining the history of the disease nearby, the size of the susceptible population, and the conditions that permit transmission.

Surveillance reports what has already happened. Forecasting systems exist for several livestock diseases but not for glanders. The problem addressed here is to build a population-level risk model for glanders from public data and to measure how well it works. It has four parts:

1. **Data assembly.** No public source records glanders occurrence completely. Occurrence must be reconstructed from the literature and joined to population and environmental covariates at a common resolution.
2. **Label construction.** The outcome is rare, so accuracy is useless and a negative observation needs an explicit meaning.
3. **Model selection under distribution shift.** A model fitted on some countries must be evaluated on countries it has never seen, because that is the deployment condition.
4. **Threshold economics.** A probability is not an alert. The operating point must be defensible against the number of warnings it produces.

3. Objectives

There is one research objective and three engineering objectives.

1. **Research.** Assemble a country-year panel for glanders from public sources not previously applied to the disease, compare six model families under three validation schemes, and report the achievable skill, including negative results.

2. **Prediction as a service.** Package the trained model as a versioned, portable artefact behind a REST interface, so any consumer can call it without knowing how it works.
3. **Cited explanation.** Build a retrieval-augmented layer that explains a prediction from official guidance, cites every claim, and refuses rather than guesses.
4. **Owner alerts.** Convert a calibrated probability into a small number of messages, using threshold, suppression and deduplication rules so that a rare disease does not train its users to ignore warnings.

All four share one requirement: each component must be independently testable, independently replaceable, and free of any hardcoded reference to glanders, so the platform can be repointed at another disease without altering the components.

4. Scope and Limitations

The study is scoped to country-year resolution. District-level modelling for India was envisaged in the project proposal but is not attempted, because the required covariates could not be assembled from public sources within the project period. The explanation and alert layers are resolution-agnostic and apply unchanged once a district-level model exists. The model is a population-level tool: it does not diagnose an individual animal and does not replace laboratory confirmation.

Three limitations shape every result that follows. The label comes from published reports, so a negative may mean no disease or no reporting. Occurrence is reconstructed by an automated extractor that has not been validated against a hand-annotated set. And with 172 positive observations the uncertainty on every headline metric is substantial, which is quantified rather than hidden.

Chapter 2

Literature Review and Gap Analysis

1. Epidemiology and Surveillance

Glanders is treated in the veterinary literature as a neglected zoonosis with high case fatality and a real occupational risk to handlers [1, 2]. A recent systematic review and meta-analysis consolidates seroprevalence estimates and identifies management practice and animal movement as the dominant drivers [3]; regional reviews document continued circulation in Brazil [4], and the disease remains endemic or repeatedly reported across parts of Africa, Asia and the Middle East.

Glanders persists where equid populations are large, animals move often, and surveillance is weak. The first two are measurable from public data. The third decides whether the disease is visible at all, and is therefore a confounder rather than a covariate.

Indian surveillance is the strongest evidence base. Malik et al. described emergence and re-emergence through outbreaks reported between 2006 and 2011 [7], Singha et al. reported serological surveillance for 2015 to 2018 [6], and the national programme covering 2019 to 2023 produced a prevalence estimate with a risk analysis [5]. These studies establish where glanders occurs and which animal-level factors the risk analysis identified. None builds a predictive model, and none uses population or reanalysis covariates as predictors.

2. Existing Forecasting Systems in India

The closest operational system is NADRES v2, the National Animal Disease Referral Expert System run by ICAR-NIVEDI [39]. It combines outbreak records from AICRP centres with meteorological and remote-sensing covariates and issues district-level forewarnings about two months ahead, using a model family that overlaps with the one used here.

This project differs in four ways, all complementary. It is glanders-specific, where a national system must spread effort across many diseases. It uses only public data, so anyone can reproduce and audit it. It attaches a cited explanation to each prediction, which matters when the reader is an owner rather than a trained analyst. And it delivers to the owner directly, with rules that suppress routine noise. NADRES has institutional access this project lacks; this project has a level of local explainability that a multi-disease platform does not attempt.

The comparison also shows that the architecture chosen here is conventional for this problem class. The open question is not whether to forecast but how to make a forecast understandable and actionable where it is used.

3. Spatial Risk and Machine Learning

Spatial work on glanders is descriptive. Fonseca-Rodríguez et al. applied spatiotemporal cluster detection to Brazilian data [8], and species distribution and ecological niche models have been compared as tools for livestock disease [9]. Cluster detection identifies where cases have already concentrated; it does not forecast which cells will be positive next period, which is the question an early warning system has to answer.

Predictive modelling of animal disease is established for avian influenza. Kane et al. compared ARIMA and random forest models for H5N1 outbreaks [11], Walsh et al. used machine learning to target active surveillance [12], and Yoo et al. and Alshatti et al. built farm-level risk models for highly pathogenic avian influenza in Korea and Kuwait [13, 10]. That literature supplies the methodological template used here: assemble occurrence, population and environmental covariates at a common unit, fit tree-based ensembles, and evaluate on held-out periods. The template has been applied to influenza, not to glanders.

The only machine learning work specific to *B. mallei* is spectroscopic identification. Moawad et al. built a Raman assay in which learned models classify spectra [14], and Stöckel et al. reported spectroscopic detection

of *B. mallei* and *B. pseudomallei* [15]. Both are laboratory diagnostic problems. Neither addresses geographic or temporal distribution, and neither produces a population risk estimate.

Using published literature as structured occurrence data has precedent in epidemic intelligence [16, 17]. This work makes sentence-level event extraction the primary source of the label rather than a supplementary signal, because official outbreak records are incomplete for many countries. The label therefore inherits the reporting biases of the literature, which is analysed rather than assumed away.

4. Prior Art Audit

To test the novelty claim rather than assert it, a reproducible set of queries was run against Europe PMC. Each row of Table 2.1 is a query count; the query strings and retrieved titles are retained with the project code.

Table 2.1: Europe PMC hit counts for combinations of glanders with each data source and method used in this study.

Query	**Hits**
glanders (baseline corpus)	6674
glanders + machine learning	75
glanders + tree ensembles	40
glanders + forecasting	139
glanders + FAOSTAT	7
glanders + NASA POWER	0
glanders + World Bank indicators	22
glanders + equine population denominator	54
glanders + Europe PMC text mining	378
glanders + chained/stacked models	36
glanders + random forest (risk context)	8
glanders + niche / distribution model	62
glanders + remote sensing / reanalysis	10
glanders + spatio-temporal clustering	463
glanders + gradient boosting	4

The decisive cells pair glanders with NASA POWER, which returns zero records, and with gradient boosting, which returns four, none of them a glanders risk model. The nearest prior work is descriptive cluster detection in Brazil and a risk-factor analysis of the Indian serosurvey.

5. Research Gap

The literature divides into three disjoint groups. Epidemiological studies describe where glanders occurs. Machine learning studies of *B. mallei* classify laboratory spectra. Veterinary risk models exist but are built for influenza. No study joins an occurrence corpus, a livestock census, development indicators and reanalysis climatology into a single panel for glanders, and none evaluates a calibrated predictive model on that panel.

A second gap appears once the question moves from analysis to use. A probability on a dashboard does not reach the person who owns the animals, and a risk score without a stated basis is not something a field officer will act on. Both bodies of literature stop short of the delivery problem: how a forecast becomes a cited explanation, and how an explanation becomes a message an owner receives and trusts.

The contribution therefore has two parts. The empirical part is a calibrated population-level risk model for glanders, built on data sources not previously applied to the disease and evaluated under a forward temporal split and a country-holdout split. The architectural part is a decomposition of the delivery problem into three components with fixed contracts, so that prediction, explanation and alert can be built, tested and replaced independently, and so the same structure carries over to another disease.

Chapter 3

System Architecture

1. What the Platform Has to Do

The platform answers one question and produces three outputs from it. The question is whether glanders is likely to be documented in an area next period. The outputs are a number on a map, an explanation in words, and a message to an owner. Each is read by a different person, at a different time, through a different device, and each can fail without the others failing.

That asymmetry is why the system is three components rather than one application. A district officer needs the prediction immediately and the explanation on demand. An owner with a feature phone needs a short message and nothing else. A researcher auditing a decision needs the cited passages months later. One service serving all three would tie alert delivery to the latency of an LLM call, and would make the risk model impossible to test on its own.

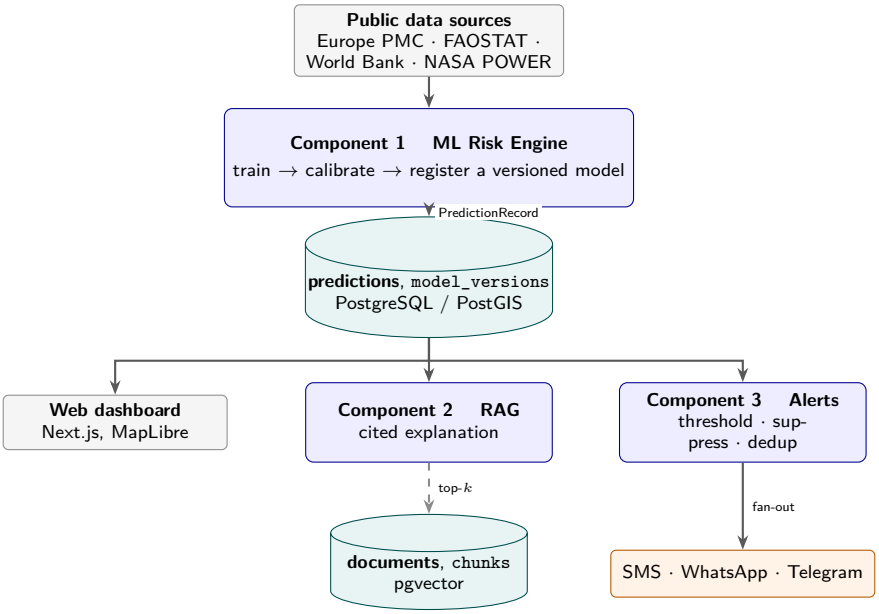


Figure 3.1: The platform decomposed into three components. The ML risk engine writes one `PredictionRecord` per area and period; the dashboard, the RAG layer and the alert service consume it independently. No component reads another component's tables, so any one of them can be replaced without touching the other two.

2. Design Principles

Five rules govern the decomposition. Every later design decision follows from them, and they are the part of the project that transfers to other diseases.

One responsibility per component. The risk engine turns features into a calibrated probability, the RAG layer turns a prediction into a cited explanation, and the alert service turns a prediction into a delivered message. None computes what another could compute instead.

Contracts before implementations. Components exchange versioned schemas, never shared tables or in-process calls. The contract is defined first, so each side can be built and tested against a fixture. This is what let the RAG layer be developed while the alert service was still being written.

The core is disease-agnostic. Nothing hardcodes glanders. A disease is a configuration record: which label extractor, which knowledge-base collection, which templates, which thresholds. A second disease is a configuration change plus a retrain, not a code fork.

Fail soft, never fail silent. A degraded component returns a degraded answer. If the RAG layer is unavailable the alert still goes out, with the level and standing action text but no explanation paragraph. If retrieval finds nothing above the relevance floor, the layer returns a refusal plus the model's own SHAP factors. Neither case produces a wrong statement, and neither produces silence.

Every output is traceable. Each prediction carries the model version that produced it, each explanation the document identifiers it was grounded in, each alert the prediction that triggered it and the rule outcomes that allowed it. A decision can be reconstructed after the fact, which is the minimum requirement for a tool that advises on animal disease.

3. The Three Components

Table 3.1: Component responsibilities, interfaces and failure behaviour.

Component	Responsibility	Interface	On failure
ML risk engine	Features to calibrated probability, with attribution	`POST /v1/predict`, `GET /v1/risk-map`	`503` with the champion version id; the last good prediction stays in the table, marked stale.
RAG layer	Prediction plus question to a cited explanation	`POST /v1/explain`, `POST /v1/chat`	Alert text is sent without the explanation; chat returns a refusal rather than an ungrounded answer.
Alert service	Prediction to a delivered, deduplicated message	`POST /v1/alerts/evaluate`; consumes `predictions_ready`	Alerts queue in Redis and drain on recovery; map and explanations are unaffected.

4. Component Contracts

A contract is a JSON schema plus the guarantees that accompany it. A narrow contract is what makes a component replaceable, so all three are deliberately small.

The risk engine emits a `PredictionRecord` containing the area, period, calibrated probability, discrete risk level, model identifier, a data-freshness block and the top SHAP factors. Three guarantees accompany it: the probability is calibrated on held-out data and may be read as a frequency; the risk level is derived from the probability by the threshold configuration, so the two cannot contradict each other; and the top factors are computed from the model itself, so they remain available when the RAG layer does not.

The RAG layer accepts an `ExplainRequest`, which wraps a prediction with a question and a language, and returns an `ExplainResponse` carrying an answer, a list of resolved citations, a `grounded` flag and a refusal reason. A refusal carries no answer text, so a caller cannot render one as though it were an explanation.

The alert service emits an `AlertEvent` carrying the alert identifier, area, period, current and previous level, the prediction it derives from, the outcome of each rule, and a recipient count. The identifier is derived deterministically from its key, for reasons given in Chapter 7.

The components share a database but not tables: each table has one owning writer and the others read it, and no component updates another's rows. Spatial columns use PostGIS and document embeddings use pgvector, so spatial and vector queries run beside the relational joins rather than in a second store that would need synchronising.

5. Modularity and Reuse

The decomposition justifies its overhead only if the parts are reusable. Five paths follow, each a small change rather than a rewrite.

A second disease supplies a new gazetteer and lexicon for extraction, a new guidance collection for retrieval, and its own tuned thresholds. The modelling code does not change, because it never refers to glanders by name.

A second LLM provider is a configuration value. The RAG layer calls a provider interface with one completion method, and an adapter supplies it for either a local Ollama model or a hosted API. The retrieval and guardrail code is provider-independent, so citation checking and refusal behave identically either way.

A second notification channel is an adapter plus a template. The adapter exposes one send operation returning a delivery status, and the rules engine never learns which channel will carry the message.

The trained pipeline is a scikit-learn compatible object [32] exported with a manifest recording feature order, calibration and metrics. Any Python process can load it without the API, which is what makes offline batch scoring and notebook analysis possible.

A replacement model needs no downstream change, since the record carries `model_id` and the registry keeps previous versions, so a candidate can be scored against the incumbent on a shadow period before promotion.

6. Deployment and Status

The services are containerised in one Docker Compose file: API, worker, PostgreSQL with PostGIS and pgvector, Redis, and the frontend. Local demonstration and server deployment therefore have the same structure. A Celery beat job runs the feature job, then the prediction job, then publishes `predictions_ready`; explanations and alerts are computed from that event, so a failed explanation never blocks a prediction and a failed prediction never sends an alert about stale data.

The report distinguishes throughout between what has been built and what has been specified. The risk engine is built and evaluated, and every number in Chapter 8 is reproducible from the committed scripts. Model serving, the RAG layer and the alert service are specified with fixed contracts, exact rule semantics and stated evaluation protocols, but they are not deployed, so no measured performance is reported for them.

Chapter 4

Methodology

1. Data Sources

Four publicly accessible sources are used: a literature index, a livestock census, a development indicator set and a reanalysis climate archive [38, 35, 37, 36]. None had previously been applied to glanders risk modelling, and each contributes a distinct kind of evidence.

Table 4.1: Data sources and their role in the panel.

Source	Role	Content	Coverage
Europe PMC	Animal history	Occurrence events extracted from abstracts	6,674 records, 263 events
FAOSTAT	Animal population	Stocks of horses, asses, mules, camels, cattle, sheep and goats	68,449 records, 1961–2023
World Bank	Socio-economic context	Ten development indicators and country centroids	217 countries, 2000–2024
NASA POWER	Environment	Temperature, precipitation, humidity and wind	Six measures, 211 centroids, 2005–2024

Extracted events rise in volume after 2015, which reflects growth in open access veterinary publishing as much as any change in incidence. This is one reason the label is treated as documented activity throughout.

2. Panel Construction

All four sources are keyed on ISO 3166 alpha-3 country code and calendar year, so they join without ambiguity once names are reconciled. FAOSTAT is streamed from its 32 MB compressed bulk archive and filtered to the seven stock series of interest rather than loaded whole.

Country names are matched to ISO-3 through `pycountry` with explicit overrides for non-standard forms such as “Türkiye”, “Viet Nam” and “Côte d’Ivoire”. World Bank regional and global aggregates are excluded by requiring a three-character ISO code, which stops rows such as “World” or “South Asia” from entering the panel as though they were countries.

The full grid holds 6,700 country-year rows across 268 countries for 2000 to 2024. Three restrictions produce the modelling matrix: only equine-holding countries are retained, defined as a maximum equid stock of at least one thousand animals; only 2005 to 2023 is used; and only rows with an observable next-year label are kept. The result has 3,078 rows, 44 features and 172 positive labels.

Extracted events are de-duplicated on country, year and source record identifier, so one paper reporting several cases cannot contribute multiple counts. Missing values are median-imputed inside the training pipeline, and each affected column keeps a missing-indicator feature so the test period cannot influence imputed values. Missingness is low: 0.00% for FAOSTAT, 1.68% for the World Bank, 2.19% for climate, and 0.99% overall, with no column above 5.8%. A country-year is labelled negative only after the country passes the equid-stock filter, so a blank is never read as proof of absence.

3. Label Definition

The target is binary: whether a glanders occurrence was documented in country i during year $t+1$. Occurrence is extracted from Europe PMC by matching, within one sentence, a country name from a curated gazetteer

and a term from an occurrence lexicon containing *outbreak*, *case*, *detected*, *seropositive* and *surveillance*. An event is recorded for the year stated in the sentence, or for the publication year where no year is stated.

Because the label depends on publication it is termed *documented activity*. A negative may mean no disease or no report, so the model predicts surveillance-observable activity rather than biological presence.

4. Feature Engineering

Four groups are built, all from information available at year t or earlier. Seven disease-history features cover prior-year occurrence, a five-year running total, years since the last event, and world and regional activity shares. Fourteen population features cover log equid and livestock stocks, equid share, density, growth and the seven species counts. Ten climate features cover temperature, precipitation, humidity and wind, each paired with an anomaly against the country's own expanding mean, which separates unusual conditions from latitude. Eleven socio-geographic features cover income, agriculture, GDP, rural share, connectivity and population, and largely proxy surveillance capacity.

Three derived features need definition. Equid share is horses, asses and mules divided by total farm animals. Livestock units apply metabolic weights of 0.8 for cattle and camels and 0.1 for small ruminants, so one number expresses grazing pressure across species. Years since the last event uses 99 to denote no event in the panel window.

Selection ran in two stages. Columns were first removed for having no measurable association: mutual information below 0.002 and absolute Spearman correlation below 0.02. Redundancy was then addressed by dropping one member of each pair correlated at or above 0.90, and by dropping a column that was exactly the sum of three others, which a tree can construct itself. No projection or principal component analysis was applied: with 44 features and 3,078 rows the binding constraint is signal rather than dimensionality, and the ablation in Section 4. shows that reduced sets transfer better than the full set.

5. Models

Six families are compared, spanning additive to fully interactive. Logistic regression is the additive baseline with balanced class weights. A single shallow decision tree is included because a non-specialist can read it, and because the gap between one tree and an ensemble is itself informative. Random forest and extra trees are bagged ensembles [20], differing in split randomisation; XGBoost and LightGBM are boosted ensembles differing in tree growth [18, 19]. Tuning used grouped cross-validation on country. Neural networks were excluded on evidential grounds: with 172 positives and 44 features a network has far more parameters than examples, and would memorise the training period.

6. Evaluation Design

Three schemes are reported, each answering a different question. A temporal holdout trains on 2005 to 2017 and tests on 2018 to 2023, with the test period never observed during fitting; it answers whether the model can be used in following years. Leave-country-out uses five folds grouped by country, so every prediction is for a country absent from training; it answers whether the model transfers, and it governs model selection because it is the harsher test and the closer approximation of deployment. Repeated stratified cross-validation, five folds repeated three times, is a stability check on the ranking metrics.

Precision-recall AUC is the primary metric. At 5.59% prevalence ROC-AUC is dominated by the true-negative mass and overstates performance [22, 23], so it is reported only for comparability. Brier score and expected calibration error are reported because an uncalibrated probability cannot support a threshold-based alert [25, 26]. Precision, recall, specificity, F1 and the Matthews correlation are computed at a threshold tuned on training data only, and bootstrap percentile intervals with one thousand resamples quantify uncertainty.

Grouped rather than random splitting is essential, because a random row split would place the same country in both partitions and let the model memorise country identity [24].

7. Leakage Controls

Six controls are applied: features use year t or earlier and only the target uses $t + 1$; the split is chronological, so no future observation enters training; hyperparameters are tuned inside training folds; calibration is fitted on training folds only; the alert threshold is fixed on training and applied unchanged to test; and imputation values and missing indicators are generated inside the pipeline.

Chapter 5

MLOps and Model Serving

1. From a Notebook to a Service

A trained model is not a component until it can be called, versioned, monitored and replaced. Three things must be true first. The model must be a portable artefact, so the object scored in evaluation is the object served, with nothing re-fitted in between. Every prediction must name the version that produced it, so a decision can be explained months later. And the service must degrade predictably, so a bad release or a missing feature does not produce confident nonsense.

The project is a sequence of numbered scripts, each producing artefacts consumed by the next, with every intermediate output cached so any stage can be re-run without repeating the stages before it. This is also what makes the pipeline inspectable: a reader can run one stage and check it before trusting the next.

2. The Model Bundle

The unit of deployment is a directory, not a pickled object, because a pickle alone cannot say what its inputs mean.

```
gw-xgb-2026.11.1/
  model.joblib        # fitted pipeline, predict_proba
  manifest.json       # feature order, dtypes, calibration, metrics
  thresholds.json     # level boundaries and hysteresis margin
  SHA256SUMS
```

The manifest is what makes the bundle reusable. It records the ordered feature list with dtypes, so serving can reject a request with the wrong columns before any inference happens. It records the calibration method and achieved error, so a consumer knows whether to read the output as a probability or only as a rank. It records the evaluation metrics, so the number in the report and the number in the registry are the same number. Because the bundle is a directory with a manifest, the API, a batch job and a notebook load it with identical behaviour.

Every training run also writes a record containing the git commit, data version, feature list and hash, hyperparameters and the full metric set for all three schemes. Promotion is decided on leave-country-out precision-recall AUC, the metric that governs selection, rather than on the more flattering temporal holdout. A candidate becomes champion only if it beats the incumbent by more than the bootstrap interval width, its expected calibration error is below the deployment ceiling of 0.05, its feature schema matches what serving can supply, checked by name and order, and it has scored a shadow period without a prediction outside the plausible range. The previous champion is retained, so promotion is reversible, and a bundle failing its checksum or schema check is replaced by the previous champion instead of being served.

3. Serving Design

The interface is deliberately coarse: a caller asks for a prediction for an area and period and receives the full record, rather than assembling one from several calls.

Every response carries the `model_id` that produced it and a `data_freshness` block. The alert service reads that block before sending anything, so a prediction computed on stale upstream data cannot produce a message.

The order of operations in Figure 5.1 is a reliability choice. Validation runs first, so a malformed request costs a schema check rather than a model load. The cache is consulted before inference because risk maps are read

Table 5.1: HTTP interface of the risk engine.

Endpoint	Method	Purpose and contract
`/v1/predict`	POST	Score one area and period from a supplied or stored feature vector. Returns one `PredictionRecord`.
`/v1/predict/batch`	POST	Score up to 500 area-period pairs, with per-item errors rather than failing the request.
`/v1/risk-map`	GET	GeoJSON of the latest predictions for a period, filtered by risk level. Serves the dashboard directly.
`/v1/models`	GET	Registry listing: model id, promotion date, metrics, feature-schema hash.
`/v1/models/{id}`	GET	The manifest of one version, for inspection by a consumer.
`/v1/health`	GET	Liveness and readiness, where readiness checks that a champion bundle is loaded and its schema resolves.

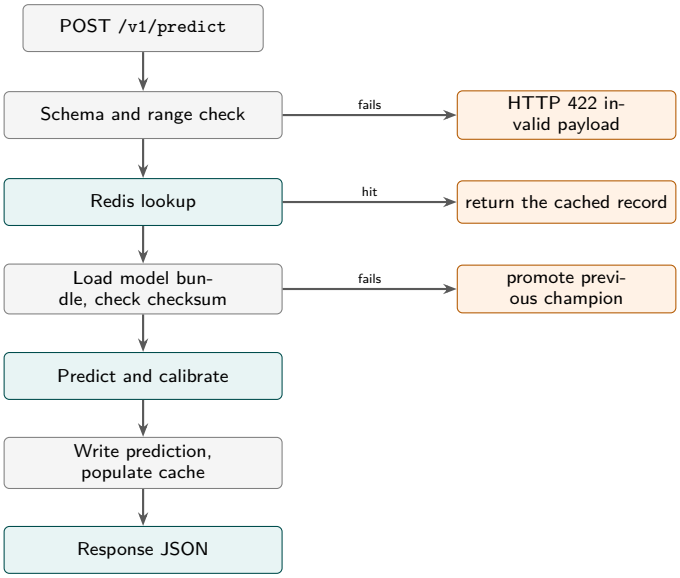


Figure 5.1: Request lifecycle of the prediction endpoint. Validation precedes model loading, the cache is consulted before inference, and a bundle that fails its checksum or its feature-schema check is replaced by the previous champion rather than served.

far more often than recomputed, and the same area-period is requested repeatedly in one dashboard session. Inference is deterministic for a fixed bundle, so a cache key of model id, area and period is safe, and the model id is part of the key rather than a version counter so promotion invalidates the cache. Responses are written back to the predictions table as well as cached, so the map does not depend on a warm cache.

The endpoint is specified against a budget rather than measured, since the service is not deployed. The budget is set so a full risk map for one period is interactive.

Table 5.2: Target latency for a single prediction request.

Stage	Target	Note
Schema and range validation	2 ms	Rejects malformed input before any I/O
Cache lookup	3 ms	Redis, keyed on model, area, period
Feature retrieval	25 ms	Indexed PostGIS and relational joins
Model inference	8 ms	A single row through a tree ensemble
Calibration and record assembly	2 ms	Isotonic lookup plus rounding
Total, cache miss	40 ms	
Total, cache hit	4 ms	

The explanation endpoint has a different budget of two seconds, because it includes generation. That is why the alert service never calls it inline: a message must not wait on an LLM.

4. Monitoring and Retraining

Three classes of signal are monitored. Service health covers request rate, error rate and latency percentiles per endpoint, exposed on a Prometheus endpoint and charted in Grafana; this catches the service being down, which is the least interesting failure. Data drift compares each feature's current distribution with the training distribution using a population stability index and a Kolmogorov-Smirnov statistic, and reports any feature outside its band, because a tree asked to extrapolate outside its training range returns a confident answer with no basis. Prediction drift and calibration decay track the mean predicted probability per period against the observed event rate as labels arrive: a model that keeps saying 0.2 while the observed rate falls to 0.05 has lost calibration even with its ranking intact.

Retraining follows from the third signal. Recalibrate when the observed-to-predicted ratio leaves a band for two consecutive periods, and retrain when a new year of labels is available and the champion's leave-country-out score has degraded beyond the bootstrap interval.

The pipeline runs linting, unit tests and a smoke test on every push. Unit tests cover what can be checked without data: schema validation, the level-threshold function including hysteresis at the boundary, idempotency key derivation, and the citation parser. A further test loads the committed bundle, scores a fixed fixture and asserts the output to four decimals. This catches the most common silent failure here, where a dependency upgrade changes numerical behaviour and every downstream number moves unnoticed.

The full computation runs in about 4.5 minutes on an Apple M4 processor with 10 cores and 16 GB of unified memory, without a graphics accelerator. Per-family fitting time at the country-year level ranges from 1.9 seconds for the decision tree to 50.7 seconds for LightGBM. One module-level seed governs fitting, cross-validation, bootstrap resampling and SHAP sampling, with software versions Python 3.12.13, scikit-learn 1.9.1, XGBoost 3.4.1, LightGBM 4.7.0, SHAP 0.54.1, NumPy 2.5.4, pandas 3.1.1, SciPy 1.18.1. Raw payloads are cached and the derived panel is written to `data/processed`, so every reported number can be regenerated without re-downloading. One caveat applies: tree ensembles are seeded but not bitwise reproducible across thread counts, so the bundle smoke test runs single-threaded.

Chapter 6

The RAG Explanation Layer

1. The Problem This Layer Solves

A calibrated probability of 0.24 is actionable for a veterinary officer and meaningless to an equine owner. The owner needs to know what happened, what it means for their animals, and what to do next. The officer needs to know which factors drove the number, because a forecast they cannot interrogate is a forecast they will not act on.

Asking a large language model is the obvious answer and the obvious hazard. A model asked about glanders answers from training data that may be outdated, may conflate glanders with melioidosis, and may invent a treatment protocol. For a zoonosis under statutory control, an invented instruction is worse than no instruction. Retrieval-augmented generation [27] changes what the model may use: it is not asked what it knows, but given a small set of passages from official guidance plus the prediction's own numbers, and required to answer only from those. Each claim traces to a passage, and the passage is shown to the reader.

The goal is therefore narrow and testable: an answer a veterinary officer would accept as consistent with the National Action Plan, and a refusal rather than a guess when the corpus does not support one.

Fine-tuning would produce fluent text with no citations and no way to tell guidance from the base model's priors, and it would freeze the knowledge at training time, so a revised plan would require retraining. Retrieval keeps knowledge in a corpus that a re-ingestion can update, makes every answer auditable by construction, and allows a refusal path that fine-tuning does not. The cost is latency and possible retrieval failure, both handled explicitly below.

2. Knowledge Base

The corpus is restricted to sources with authority, because the layer's value is that it cites them. Each document carries an authority tier, used as a tie-breaker during retrieval.

Table 6.1: Knowledge base, by authority tier.

Source	Tier	Use
DAHD National Action Plan for Glanders, 2025 [34]	1	Case definitions, surveillance obligations, movement control, quarantine, notification duties.
WOAH and WHO disease information sheets	1	Disease biology, transmission, zoonotic risk, official status.
ICAR and national laboratory guidance	2	Diagnostic sampling, laboratory referral.
Open-access research papers	3	Epidemiology, risk factors, regional situation.

Tier 3 may inform an explanation but never overrides tier 1 on a control instruction. When passages conflict the higher tier wins and the conflict is logged, because a silent contradiction between sources is exactly the failure this layer exists to prevent.

3. Ingestion, Embedding and Retrieval

Ingestion runs offline as a job. A document is parsed, split at section boundaries, and each chunk is stored with enough metadata to be cited and filtered.

Chunking is section-aware rather than fixed-width, because a fixed window cuts control instructions in half and lets the model complete a truncated sentence. The splitter respects headings first, then paragraphs, and falls back to a sliding window only for a section longer than the ceiling. Chunks target 400 to 700 tokens with 15 per cent overlap inside a section, sections under 80 tokens merge into the parent, and the retained metadata is document id, title, full section path, page, year and authority tier. The section path is kept in full, for example `NAP-2025 > 5.2 Movement control`, so a citation points to a subsection rather than to a 90-page document. Chunks are embedded with a compact sentence-transformer [28] producing 384-dimensional vectors, sufficient for a corpus this size and cheap enough to embed on a laptop; a multilingual variant is used for Hindi, because an English-only embedding degrades badly on Devanagari. Vectors live in pgvector beside the relational tables, so re-ingestion and prediction share one transaction boundary and there is no second store to synchronise. The index is HNSW with cosine distance [29]; an approximate index is acceptable because the reranker corrects recall at the head of the list, and it keeps query cost flat as the corpus grows.

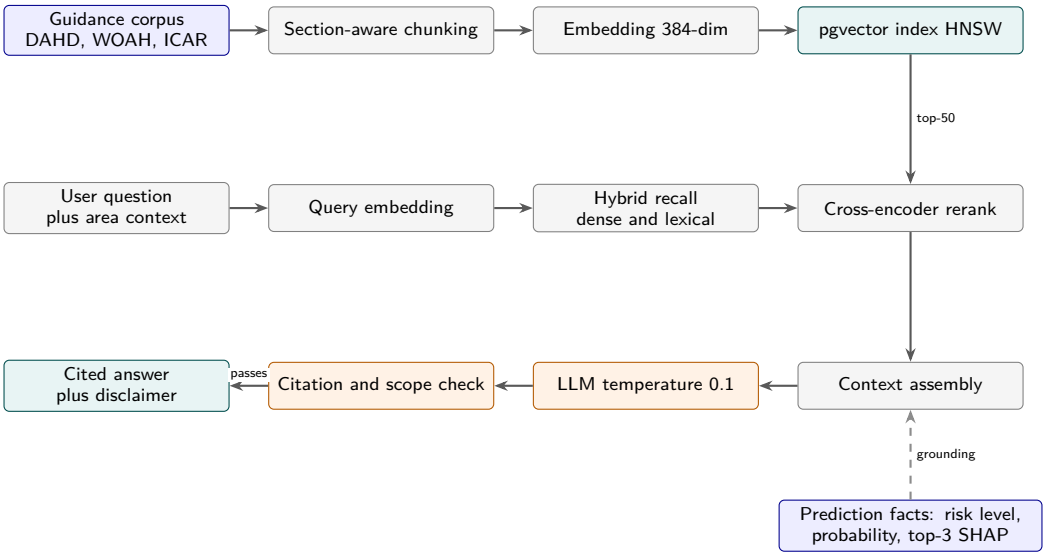


Figure 6.1: The RAG layer. Ingestion (top row) is an offline job; retrieval, assembly and generation (below) run per request. Predicted risk and SHAP factors are injected as structured facts, so the answer explains this forecast rather than glanders in general. The citation check sits after generation and can return a refusal instead of an unsupported answer.

Retrieval is hybrid, because dense and lexical search [30] fail differently here. A dense query for "what should I do if I suspect glanders" retrieves passages about suspicion, which helps. A dense query for "movement control" may miss a passage using "restriction of equid movement", where a lexical match succeeds; conversely a lexical query for a rare acronym retrieves noise a dense query ignores. Both lists are retrieved and fused by reciprocal rank fusion [31], which combines rankings without requiring comparable scores:

$$\mathrm{RRF}(d) = \sum_{i \in \{\text{dense},\text{lexical}\}} \frac{1}{k + \mathrm{rank}_i(d)}, \qquad k = 60$$

The fused top fifty are reranked by a cross-encoder scoring each passage against the question jointly. A cross-encoder is far slower per pair than an embedding dot product, which is why it runs on fifty candidates rather than on the corpus, and the top five become the context.

Table 6.2: Retrieval parameters.

Stage	Setting	Reason
Dense recall	top 30, cosine	Semantic paraphrase coverage
Lexical recall	top 30, BM25	Exact terms, acronyms, section titles
Fusion	reciprocal rank, $k = 60$	Scale-free combination of two rankings
Rerank	cross-encoder, top 50 to 5	Precision at the head of the list
Relevance floor	rerank score below 0.30 discarded	Prevents padding the context with weak passages

The relevance floor drives the refusal path. If fewer than two passages clear it, the layer returns a refusal with the SHAP factors instead of an answer, because a context assembled from marginal passages produces a confident answer built on nothing.

4. Context, Generation and Guardrails

The prompt has three parts in fixed order: a system instruction, the retrieved passages with identifiers, and the structured prediction facts. Those facts are what make the answer specific to this forecast rather than to glanders in general, and they are injected as structured data rather than prose so the model cannot restate an association as a cause. The system instruction is not influenced by user input and states four constraints: answer only from the passages, cite every claim with a passage identifier, never diagnose an individual animal, and state that the output is a population-level forecast.

Generation uses temperature 0.1 and a bounded output length. Low temperature is chosen for consistency rather than caution: two officers asking the same question about the same area should receive the same instruction. The provider sits behind a single-method interface, so the same retrieval, grounding and citation code runs against a local Ollama model or a hosted API, and switching model cannot silently disable the guardrails.

The layer is designed to refuse, and four checks can terminate a response.

Table 6.3: Guardrails and their failure behaviour.

Check	What it verifies	On failure
Retrieval floor	At least two passages above the relevance threshold	Refusal with the top SHAP factors, which need no retrieval
Citation resolution	Every cited identifier exists in the retrieved set	Answer rejected and regenerated once, then refusal
Diagnosis scope	No individual-animal diagnosis or treatment prescription	Response rewritten without the clinical claim, and the attempt logged
Out-of-scope	Question is about glanders control, not unrelated topics	Refusal stating what the layer covers

The citation check is mechanical rather than prompted. The answer is parsed for identifiers and each is looked up in the retrieved set; an identifier that does not resolve means the model cited a document it was not given, which is a hallucination whether or not the sentence happens to be true. Prompting a model to be accurate is not a control; checking its output is.

5. Evaluation

Four quantities are measured, none of them whether the prose reads well. Retrieval quality is recall@5 and mean reciprocal rank on a hand-built question set, which isolates retrieval from generation, since no generation quality helps if the right passage is never retrieved. Faithfulness is whether each claim is supported by the retrieved context, scored by a judge model and reviewed by hand on a sample; this catches a model that cites a passage but says something it does not say. Citation precision is the fraction of citations resolving to a passage that supports the sentence, which is distinct from resolution, since resolution only checks existence. Refusal behaviour is the refusal rate on out-of-scope questions and the false refusal rate in scope: a layer that never refuses is unsafe and one that always refuses is useless, so both are reported.

6. Worked Example

The intended shape of a response for a high-risk area, with citations to the corpus above.

> **Question.** Why is district IN-UP-057 rated HIGH this month?
>
> **Answer.** Risk is elevated mainly because the district has a high equine density and there has been reported glanders activity in surrounding districts within the last two years. Animals in areas with recent reported activity are at higher risk, and movement of equids is a recognised route of spread [P1]. Owners are advised to avoid unnecessary movement or mixing of equids, to isolate animals showing nasal discharge or skin lesions, and to inform the local veterinary authority rather than attempt treatment themselves [P1, P2]. Glanders is zoonotic, so handlers should use gloves when examining a suspected animal [P2].
>
> This is a population-level forecast for the area. It is not a diagnosis of any individual animal, and suspected cases require laboratory confirmation [P1].

Every sentence is attributable to a retrieved passage: the control instruction cites the National Action Plan, the zoonotic warning cites the WOAH sheet, and the disclaimer is appended by the layer rather than generated.

Because the layer is specified rather than deployed, the validation that would precede deployment is stated so the claim is falsifiable: recall@5 above 0.85, faithfulness above 0.90 on a manually reviewed sample of 100 answers, citation precision above 0.95 with every failure inspected, zero individual-animal diagnoses across an adversarial probe set, refusal on every out-of-scope question with a false refusal rate below 10 per cent in scope, and p95 latency under two seconds on the reference hardware. These are targets, not results.

Chapter 7

The Alert Service

1. Why the Alert Is the Hard Part

The risk engine produces a probability and the RAG layer explains it, but neither reaches the person who owns the animals. For most Indian equid owners the interaction is a message on a phone, competing with everything else on that phone.

The design problem is therefore not how to send a message but how to send few enough that the ones which arrive are read. A system that warns about every area every month trains its users to ignore it, and then fails silently at the moment it was built for. Alert fatigue is the dominant failure mode of an early warning system for a rare disease, and the rules below exist to prevent it.

The service is event-driven. When the prediction job completes it publishes `predictions_ready`, and the alert evaluator consumes it, so an alert can never be sent for a period whose predictions were not recomputed. For each area with an active subscription the evaluator reads the new prediction and the area's previous alert state, and those two values are the entire input to the rules. It reads no features and no model internals, so a model change can affect alert behaviour only through the probability it produces.

2. Rule 1: Threshold with Hysteresis

The calibrated probability maps to three levels. The boundaries are configuration values derived from the operating curve in Chapter 8, not defaults and not 0.5.

Table 7.1: Level thresholds and their basis.

Level	Band	Meaning for the recipient	Basis
LOW	below the medium boundary	No message. Routine monitoring.	Below the budget where precision collapses
MEDIUM	medium to high boundary	Prepare; check animals; note the veterinary contact	Operating point where recall exceeds 80 per cent
HIGH	above the high boundary	Act: restrict movement, isolate suspects, inform the authority	Top of the operating curve, precision above 0.5

Hysteresis addresses a specific defect. With a single boundary, an area whose probability oscillates around the threshold produces an alert every period, because each crossing is a state change. Two boundaries are therefore used: an area enters HIGH at the high threshold and leaves only when the probability falls below that threshold minus a margin.

$$\text{enter HIGH at } p \geq T_H, \qquad \text{leave HIGH only at } p < T_H - \delta, \qquad \delta = 0.05$$

The margin is configurable per disease. A larger margin gives a more stable state at the cost of slower de-escalation, which suits a disease where a missed warning costs more than a late all-clear.

3. Rule 2: Suppression

Suppression decides whether a message is warranted at all, given that the level has changed. Six conditions apply and any one can stop the message.

Table 7.2: Suppression conditions.

Condition	Rule	Rationale
State change	Send only on a transition, not while the level is held	A monthly reminder at the same level carries no new information
Cooldown	No repeat of the same level within a configurable window	Prevents oscillation at a boundary from producing repeated messages
Escalation bypass	An upward transition ignores the cooldown	A worsening must never be delayed by a timer
Per-user daily cap	At most one message per recipient per day	Protects against several areas escalating at once
Quiet hours	Nothing sent between 21:00 and 07:00 local; queued instead	A warning that wakes a household at night is resented, and is not actionable then
Stale-data guard	No send if `data_freshness.passes` is false	An alert about a prediction computed on stale inputs is worse than no alert

A manual mute per area is also supported, so an officer who knows a zone is under active surveillance can silence it without disabling the service. De-escalation is handled separately: when a level falls, an optional short all-clear is sent, because a recipient told to act needs to be told when to stop. Whether it is sent is a configuration choice, since an all-clear costs a message but prevents a user from maintaining a restriction indefinitely.

4. Rule 3: Deduplication

The service must be safe to re-run. Jobs are retried, deployments restart workers, and a model may be re-scored for the same period; none may produce a second message. Deduplication is enforced at three levels, each with its own key and database constraint.

Table 7.3: Deduplication levels.

Level	Key	What it prevents
Alert	`(area_id, period, level)`	A re-run of the job creating a second alert for the same period
Recipient	`(alert_id, subscription_id)`	One subscriber receiving the same alert through two fan-out passes
Send	`(alert_id, subscription_id, attempt)`	A retry after a timeout delivering a message that had in fact succeeded

The first two are unique database constraints, so protection does not depend on application logic being correct. The alert identifier is derived deterministically from its key rather than generated at random, so two concurrent workers computing the same alert arrive at the same identifier and one loses the insert; a random identifier would let both succeed. The send-level check reads the delivery log before each attempt, which handles the case where a provider accepts a message, the response times out, and the retry would otherwise deliver it twice.

5. Fan-Out, Channels and Delivery

Once an alert record exists, one delivery job per subscription is queued through Redis. A channel adapter receives a rendered message and a recipient and returns a delivery status, with adapters for SMS, WhatsApp and Telegram, the last used for demonstration because it needs no commercial onboarding.

Retries use exponential backoff with jitter: a permanent failure such as an invalid number is not retried, a transient one such as a provider timeout is, and the distinction is recorded in the delivery log, because a silent permanent failure leaves a subscriber who believes they are covered and is not. Templates are separate from the rules engine, so the renderer produces the words in the recipient's language from a level and a set of facts; adding a language adds a template set rather than touching the rules.

A HIGH alert as rendered for SMS. Length is bounded because a carrier truncates long messages, and a truncated message that loses the action instruction is worse than a short one.

```
GLANDERS WATCH  |  IN-UP-057  |  NOV 2026
Risk: HIGH (was MEDIUM)
Elevated risk for your area. Avoid unnecessary
movement or mixing of equids. Isolate animals with
nasal discharge or skin lesions and inform the local
```

```
veterinary authority.
Helpline: <district veterinary contact>
Forecast for the area, not a diagnosis.
```

The same alert through a richer channel carries the explanation paragraph from the RAG layer with its citations. The rule that decides whether to send is identical; only the adapter and template differ, which is what lets one alert serve a feature phone and a smartphone without duplicating the decision logic.

6. Evaluation of the Alert Layer

The service is evaluated as a warning system, not as a messaging system. Alert precision is the fraction of HIGH area-periods followed by a documented event, and it determines whether recipients trust the message. Event coverage at budget is the operating curve already reported in Chapter 8, and the thresholds are read from it. Noise per recipient is messages per subscriber per year, which predicts whether the service will still be read in its second year. Delivery reliability is the fraction of queued deliveries reaching a terminal success state within the retry window, reported per channel. The first two come from historical predictions and labels; the last two need a deployed service and are stated as the protocol that would measure them.

Chapter 8

Results and Discussion

1. Dataset Characteristics

The modelling matrix holds 3,078 country-year observations, of which 172 are positive: a prevalence of 5.59% and an imbalance of 16.9 to one. 56 of 162 countries record at least one positive, 30 of them in two or more years, and the top five countries hold 35.5% of all positives.

Three properties of this distribution matter. The imbalance rules out accuracy as a metric and motivates precision-recall analysis. The temporal persistence, with 30 of 56 positive countries recurring, justifies historical features. The geographic concentration is the first indication that the label tracks reporting capacity, since those countries operate reference laboratories.

2. Model Comparison

Table 8.1 reports discrimination under all three schemes. The temporal holdout is the realistic test and leave-country-out is the harder one.

Table 8.1: Discrimination at the country-year level. Brackets give 95 per cent bootstrap intervals on the temporal holdout.

Model	PR-AUC (time)	95% CI on PR-AUC	ROC-AUC (time)	PR-AUC (country)	PR-AUC (rep. CV)	Brier	Prec.	Rec.
Logistic regression	0.454	[0.34, 0.58]	0.812	0.223	0.414	0.0547	0.25 / 0.57	
Decision tree	0.389	[0.29, 0.50]	0.861	0.267	0.360	0.0602	0.49 / 0.53	
Random forest	0.554	[0.43, 0.66]	0.870	0.304	0.476	0.0503	0.53 / 0.50	
Extra trees	0.530	[0.42, 0.64]	0.854	0.363	0.476	0.0535	0.47 / 0.51	
XGBoost	0.485	[0.37, 0.60]	0.844	0.281	0.423	0.0536	0.61 / 0.38	
LightGBM	0.496	[0.38, 0.61]	0.841	0.339	0.397	0.0534	0.65 / 0.30	

Against a base rate of 5.59%, the strongest temporal result of 0.554 is an approximately 10-fold lift. Every tree-based model beats the logistic baseline, and the single decision tree is weakest of the six, which indicates that the relationships are interactive rather than additive.

The more important pattern is the systematic decline under leave-country-out validation. For Random Forest the score falls from 0.554 to 0.304, and the decline exceeds the bootstrap interval width for every family. This is the signature of learning country-specific structure, and it is the estimate to quote for deployment.

3. Calibration

Predictions from the selected model are well calibrated, with an expected calibration error of 0.016 and a Brier score of 0.0535. A stated risk of twenty-four per cent corresponds to an observed event frequency near twenty-four per cent. This is what makes a threshold-based alert defensible, since a message quoting a probability must be numerically trustworthy.

4. Ablation

Feature groups were removed and each configuration retrained, with results in Table 8.2.

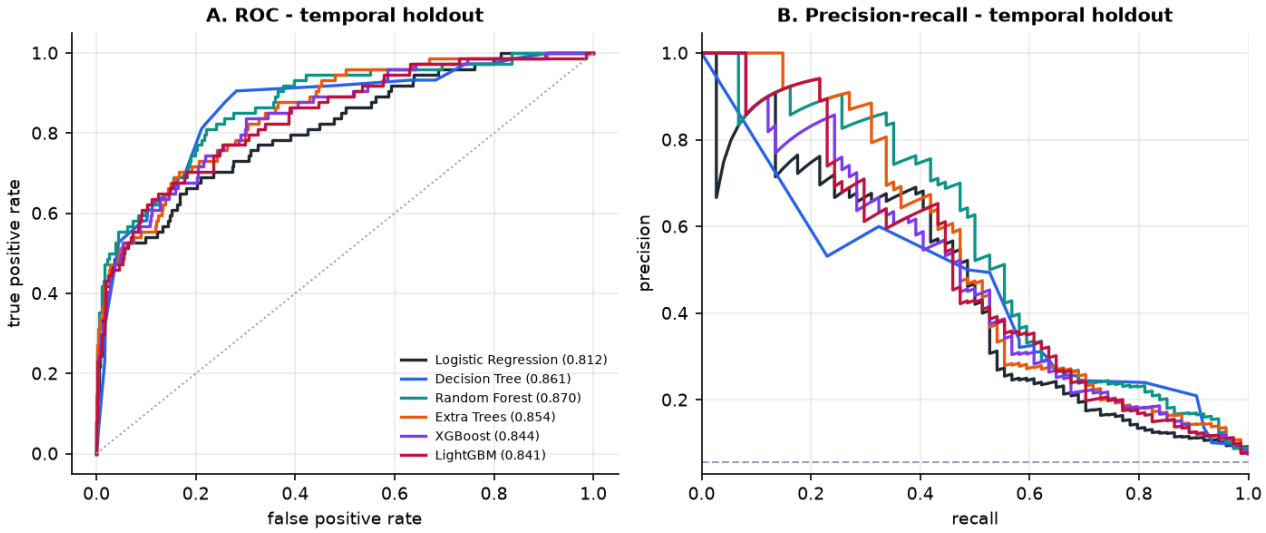


Figure 8.1: ROC curves (panel A) and precision-recall curves (panel B) on the temporal holdout. At this prevalence the ROC panels compress into a narrow band while the precision-recall panels separate the families.

Table 8.2: Ablation settings, using XGBoost with fixed hyperparameters.

Feature set	Columns	PR-AUC (time)	PR-AUC (country)
History only	7	0.407	0.174
Population only	14	0.432	0.241
Climate only	10	0.395	0.078
Socio-geography only	11	0.469	0.160
History and population	21	0.451	0.321
All groups	42	0.487	0.259
All groups plus surveillance proxy	44	0.497	0.261

Two results are more informative than the headline score. The reduced history and population set of 21 columns reaches 0.321 under leave-country-out validation, against 0.259 for the full 42 column set: added covariates improve in-distribution fit and reduce transfer, because they let the model fit country-specific peculiarities. Climate covariates alone reach only 0.078, marginally above the 5.59% floor, which indicates that annual country-level climatology largely encodes latitude and therefore country identity. A separate probe adding research-output proxies as two columns moves the leave-country-out score by +0.001, weak evidence that the model is not purely tracking publication volume, but not proof.

5. Chained Pipeline

Table 8.3 reports the effect of appending the stage-one population forecast to the occurrence model.

Table 8.3: Effect of chaining the population forecast into the occurrence model.

Variant	PR-AUC (time)	ROC-AUC (time)	PR-AUC (country)
Plain (no chained feature)	0.530	0.854	0.363
Chained (population forecast added)	0.527	0.853	0.344

The population forecast performs well as a forecast, with a mean absolute error of 0.424 on the log scale and a Spearman correlation of 0.976. The intensity stage is weaker but usable, at a mean absolute error of 0.150 and a rank correlation of 0.465 among positive observations. The hurdle product ranks expected annual burden with a Spearman correlation of 0.326, and its highest decile captures 70 of 106 observed events.

Appending the forecast changes temporal PR-AUC by -0.003, inside the noise band. The explanation is structural rather than a tuning failure: population magnitude is already a direct feature, so forecasting it and appending the forecast supplies a less accurate copy of information the model holds. Chaining should assist where the intermediate quantity cannot be observed directly, which is not the case here. The result is kept rather than removed, because it is reproducible for this problem class.

6. Alert Economics

Figure 8.2 gives the operating characteristics of the model as a warning system. Alerting the six per cent of country-years at highest predicted risk captures about half of all observed events, and alerting twenty-eight per cent captures eighty per cent. Uniform alerting would need full coverage for the same recall, so the model is roughly eight times more efficient than arbitrary selection at the fifty per cent operating point.

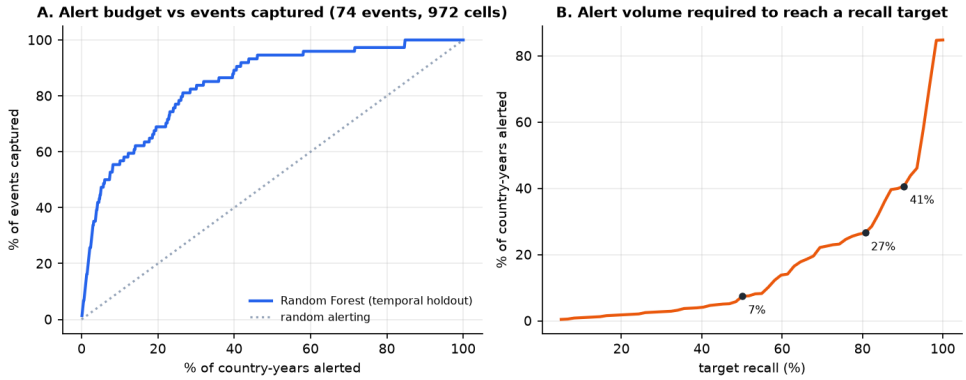


Figure 8.2: Alert economics. Panel A gives the fraction of events captured as the alert budget grows, against random alerting. Panel B gives the alert volume required to reach a target recall.

7. Subgroup Analysis

Table 8.4 reports leave-country-out performance by region. Skill is lowest in Sub-Saharan Africa and in low-income countries, which are the settings where early warning would help most. The cause is label sparsity rather than a model defect, since there are few recorded events to learn from in those strata.

Table 8.4: Leave-country-out performance by World Bank region.

Region	Observations	Events	ROC-AUC	PR-AUC
East Asia & Pacific	456	48	0.865	0.396
Europe & Central Asia	912	31	0.826	0.163
Latin America & Caribbean	551	12	0.903	0.381
MENA, Afghanistan and Pakistan	380	23	0.834	0.220
Sub-Saharan Africa	665	27	0.699	0.073

8. Discussion

The model has predictive skill. A precision-recall AUC of 0.554 against a base rate of 5.59% is not explainable by chance, it survives a strictly forward-looking temporal split, and the probabilities are calibrated to within two percentage points.

The skill does not transfer intact. Leave-country-out performance is roughly forty per cent lower in relative terms, and that is the figure for planning: any deployment to a new country should be preceded by local validation.

The principal signal is surveillance rather than biology. Three independent analyses agree. The highest-ranked features are population and connectivity measures, the label concentrates in countries with reference laboratories, and the best-scoring cells are India, the United States and Australia. Climate, which the proposal treated as a primary input, transfers poorly at this resolution.

Simpler models transfer better. The reduced feature set beats the full set across countries, and the linear baseline degrades least in absolute terms, which is a practical argument for parsimony when positive examples are limited.

Chapter 9

Conclusion and Future Work

1. Conclusion

This project set out to build a usable warning system for glanders: to establish whether population-level risk could be predicted from public data, and to design how a prediction reaches the person who owns the animals. Four sources not previously applied to the disease were assembled into a country-year panel covering 162 equine-holding countries, six model families were trained under three validation schemes, and a three-stage chained pipeline was tested. Around that model the delivery problem was decomposed into three components with fixed contracts: a risk engine producing a calibrated probability, a retrieval-augmented layer explaining it against official guidance, and an alert service turning it into a small number of messages.

Prediction is possible, and the performance is moderate. On a forward-looking temporal holdout the strongest model reaches a precision-recall AUC of 0.554 against a base rate of 5.59%, an approximately 10-fold lift, with a calibration error of 0.016. As a warning system it captures half of all observed events while alerting six per cent of country-years. The skill is fragile across geography, however: leave-country-out performance is about forty per cent lower in relative terms, and the decline concentrates where early warning would help most.

What the model learns is less encouraging. Three analyses converge: the highest-ranked features are population size and reporting capacity rather than biological drivers, the label concentrates in countries operating reference laboratories, and the best-scoring cells have the strongest publishing record. The model predicts where glanders will be documented, not where the organism is present. The chained design does not improve discrimination because the intermediate quantity is already a direct feature; that result is kept because it is reproducible.

The decomposition also proved useful: three components with versioned contracts, none hardcoding the disease, were each testable alone, so the risk engine could be evaluated before any explanation existed and the alert rules checked against a policy before any message was sent. The parts must be maintained together, and in exchange replacing the model, language model or channel is a local change.

2. Future Work

Reported activity. The most important extension is to model surveillance capacity as a quantity, as an offset on publication volume or through a two-part model estimating detection and occurrence jointly; otherwise gains in score reflect reporting rather than disease.

The label extractor. The 263 extracted events have not been checked against a hand-annotated set. Annotating a stratified sample of abstracts would give the extractor's own precision and recall, bounding label noise that currently propagates downstream.

Resolution. The proposal envisaged district-month resolution for India, where the disease is concentrated; public district covariates would need assembling, but the resolution would match the scale at which control measures are applied.

The unmeasured components. The calibration reported here is one point estimate; a production system needs drift monitoring and a retraining trigger. The RAG layer's retrieval, guardrail and protocol targets are specified but unmeasured, and running it against the corpus would give recall@5, faithfulness, citation precision and the false refusal rate.

The field. Messages per subscriber per year decides whether the service survives a second year, and only real subscribers can measure it; a district veterinary pilot would supply it. A WOAH WAHIS record set [33] as a second label source would validate the literature-derived label where the two overlap, and a cost-sensitive threshold would replace the F1-optimal one with a point matched to the programme's priorities.

Bibliography

[1] Khan I, Wieler LH, Melzer F, et al. Glanders in animals: a review on epidemiology, clinical presentation, diagnosis and countermeasures. *Transboundary and Emerging Diseases*. 2013. doi:10.1111/j.1865-1682.2012.01342.x

[2] Kanani Y, Oryan A, Derakhshandeh A. Glanders: past, present and future of a neglected zoonosis as a threat in veterinary and human health. *Veterinary Medicine and Science*. 2026. doi:10.1002/vms3.70950

[3] Purnama MTE, Firdausy LW, Fikri F, et al. Seroprevalence and associated epidemiological drivers of glanders in horses: a systematic review and meta-analysis. *Research in Veterinary Science*. 2026;106416. doi:10.1016/j.rvsc.2026.106416

[4] Mota RA, Pinheiro Júnior JW. Current status of glanders in Brazil: recent advances and challenges. *Brazilian Journal of Microbiology*. 2022. doi:10.1007/s42770-022-00814-1

[5] Singha H, Shanmugasundaram K, Jhandai P, et al. Surveillance of equine glanders in India (2019–2023): prevalence estimate, risk analysis and future strategy. *Microbial Pathogenesis*. 2026;217:108597. doi:10.1016/j.micpath.2026.108597

[6] Singha H, Shanmugasundaram K, Tripathi BN, et al. Serological surveillance and clinical investigation of glanders among indigenous equines in India from 2015 to 2018. *Transboundary and Emerging Diseases*. 2020. doi:10.1111/tbed.13475

[7] Malik P, Singha H, Khurana SK, et al. Emergence and re-emergence of glanders in India: a description of outbreaks from 2006 to 2011. *Veterinaria Italiana*. 2012. PMID:22718333

[8] Fonseca-Rodríguez O, Pinheiro Júnior JW, Mota RA. Spatiotemporal analysis of glanders in Brazil. *Journal of Equine Veterinary Science*. 2019. doi:10.1016/j.jevs.2019.03.216

[9] Hollings T, Robinson A, van Andel M, et al. Species distribution models: a comparison of statistical approaches for livestock and disease epidemics. *PLOS ONE*. 2017;12(8):e0183626. doi:10.1371/journal.pone.0183626

[10] Alshatti M, Al-Hemoud A, Othman A, et al. A risk prediction model for highly pathogenic avian influenza outbreaks in Kuwait. *Frontiers in Veterinary Science*. 2026. doi:10.3389/fvets.2026.1843962

[11] Kane MJ, Price N, Scotch M, Rabinowitz P. Comparison of ARIMA and Random Forest time series models for prediction of avian influenza H5N1 outbreaks. *BMC Bioinformatics*. 2014;15:276. doi:10.1186/1471-2105-15-276

[12] Walsh DP, Ma TF, Ip HS, Zhu J. Artificial intelligence and avian influenza: using machine learning to enhance active surveillance for avian influenza viruses. *Transboundary and Emerging Diseases*. 2019. doi:10.1111/tbed.13318

[13] Yoo DS, Song YH, Choi DW, et al. Machine learning-driven dynamic risk prediction for highly pathogenic avian influenza at poultry farms in the Republic of Korea. *Transboundary and Emerging Diseases*. 2021. doi:10.1111/tbed.14419

[14] Moawad AA, Silge A, Bocklitz T, Fischer K, Rösch P, Roesler U, Elschner MC, Popp J, Neubauer H. A machine learning-based Raman spectroscopic assay for the identification of *Burkholderia mallei* and related species. *Molecules*. 2019;24(24):4516. doi:10.3390/molecules24244516

[15] Stöckel S, Meisel S, Elschner M, Melzer F, Rösch P, Popp J. Raman spectroscopic detection and identification of *Burkholderia mallei* and *Burkholderia pseudomallei*. *Analytical and Bioanalytical Chemistry*. 2015;407(3):833–843. doi:10.1007/s00216-014-7906-5

[16] Collier N. Uncovering text mining: a survey of current work on web-based epidemic intelligence. *Global Public Health*. 2012. doi:10.1080/17441692.2012.699975

[17] Lejeune G, Brixtel R, Doucet A, et al. Multilingual event extraction for epidemic detection. *Artificial Intelligence in Medicine*. 2015. doi:10.1016/j.artmed.2015.06.005

[18] Chen T, Guestrin C. XGBoost: a scalable tree boosting system. In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. 2016:785–794. doi:10.1145/2939672.2939785

[19] Ke G, Meng Q, Finley T, et al. LightGBM: a highly efficient gradient boosting decision tree. In: *Advances in Neural Information Processing Systems 30 (NIPS)*. 2017:3146–3154.

[20] Breiman L. Random forests. *Machine Learning*. 2001. doi:10.1023/A:1010933404324

[21] Lundberg SM, Lee SI. A unified approach to interpreting model predictions. In: *Advances in Neural Information Processing Systems 30 (NIPS)*. 2017:4765–4774. arXiv:1705.07874

[22] Saito T, Rehmsmeier M. The precision-recall plot is more informative than the ROC plot when evaluating binary classifiers on imbalanced datasets. *PLOS ONE*. 2015;10(3):e0118432. doi:10.1371/journal.pone.0118432

[23] Davis J, Goadrich M. The relationship between precision-recall and ROC curves. In: *Proceedings of the 23rd International Conference on Machine Learning*. 2006:233–240. doi:10.1145/1143844.1143874

[24] Roberts DR, Bahn V, Ciuti S, et al. Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography*. 2017. doi:10.1111/ecog.02881

[25] Niculescu-Mizil A, Caruana R. Predicting good probabilities with supervised learning. In: *Proceedings of the 22nd International Conference on Machine Learning*. 2005:625–632. doi:10.1145/1102351.1102430

[26] Huang Y, Li W, Macheret F, et al. A tutorial on calibration measurements and calibration models for clinical prediction models. *Journal of the American Medical Informatics Association*. 2020;27(4):621–633. doi:10.1093/jamia/ocz228

[27] Lewis P, Perez E, Piktus A, et al. Retrieval-augmented generation for knowledge-intensive NLP tasks. In: *Advances in Neural Information Processing Systems 33 (NeurIPS)*. 2020:9459–9474. arXiv:2005.11401

[28] Reimers N, Gurevych I. Sentence-BERT: sentence embeddings using Siamese BERT-networks. In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP-IJCNLP)*. 2019:3982–3992. doi:10.18653/v1/D19-1410

[29] Malkov YA, Yashunin DA. Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs. *IEEE Transactions on Pattern Analysis and Machine Intelligence*. 2020;42(4):824–836. doi:10.

1109/TPAMI.2018.2889473

[30] Robertson S, Zaragoza H. The probabilistic relevance framework: BM25 and beyond. *Foundations and Trends in Information Retrieval*. 2009;3(4):333–389. doi:10.1561/1500000019

[31] Cormack GV, Clarke CLA, Buettcher S. Reciprocal rank fusion outperforms Condorcet and individual rank learning methods. In: *Proceedings of the 32nd International ACM SIGIR Conference*. 2009:758–759. doi:10.1145/1571941.1572114

[32] Pedregosa F, Varoquaux G, Gramfort A, et al. Scikit-learn: machine learning in Python. *Journal of Machine Learning Research*. 2011;12:2825–2830.

[33] WOAH. World Animal Health Information System (WAHIS). https:/\protect\penalty\z@{}/\protect\penalty\z@{}wahis.\protect\penalty\z@{}woah.\protect\penalty\z@{}org

[34] Department of Animal Husbandry and Dairying, Government of India. *National Action Plan for Control and Eradication of Glanders in India* (Revised 2025). New Delhi: DAHD; 2025.

[35] FAO. FAOSTAT: crops and livestock products. https:/\protect\penalty\z@{}/\protect\penalty\z@{}www.\protect\penalty\z@{}fao.\protect\penalty\z@{}org/\protect\penalty\z@{}faostat

[36] NASA Langley Research Center. POWER agroclimatology API. https:/\protect\penalty\z@{}/\protect\penalty\z@{}power.\protect\penalty\z@{}larc.\protect\penalty\z@{}nasa.\protect\penalty\z@{}gov

[37] World Bank. World Development Indicators. https:/\protect\penalty\z@{}/\protect\penalty\z@{}databank.\protect\penalty\z@{}worldbank.\protect\penalty\z@{}org

[38] Europe PMC literature search API. https:/\protect\penalty\z@{}/\protect\penalty\z@{}europepmc.\protect\penalty\z@{}org

[39] ICAR-National Institute of Veterinary Epidemiology and Disease Informatics. National Animal Disease Referral Expert System (NADRES v2): livestock disease forecasting. https:/\protect\penalty\z@{}/\protect\penalty\z@{}nivedi.\protect\penalty\z@{}res.\protect\penalty\z@{}in/\protect\penalty\z@{}Nadres_\protect\penalty\z@{}v2/\protect\penalty\z@{}